

Last week: How to make more sophisticated models by engineering features

This week: How to make sure your sophisticated model predicts the variable you want well
(how to ensure your model generalizes well)

Bias-Variance Decomposition

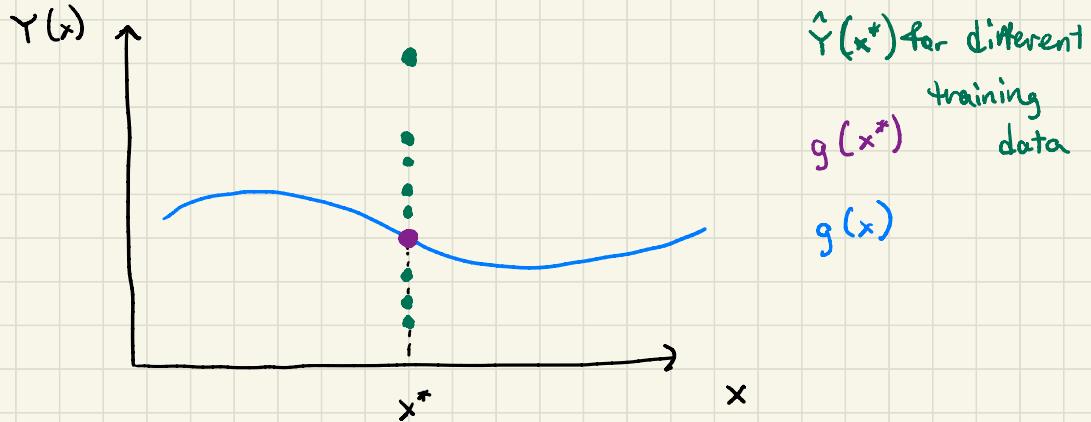
$$E[(Y - \hat{Y}(x))^2] = E[\underbrace{\hat{Y}(x) - g(x)}_{\text{model bias}^2}]^2 + E[(\hat{Y}(x) - E[\hat{Y}(x)])^2] + \sigma^2$$

we don't know g

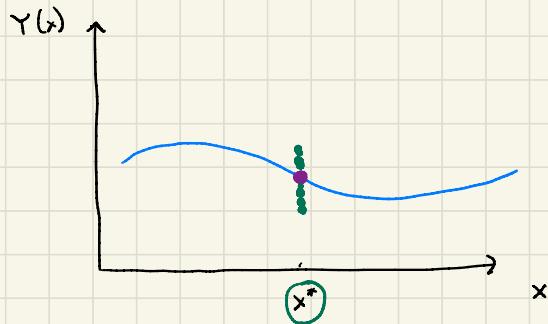
↑ ↑ ↑
model biasvarianceerror

Model Bias: How well does my model predict data it's trained on?

Model Variance: If I were to change my training data slightly, how much would the model's prediction change?

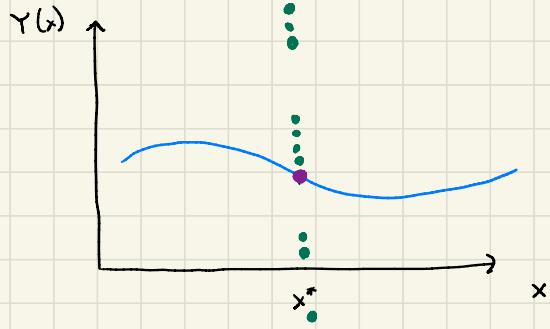


Low Bias



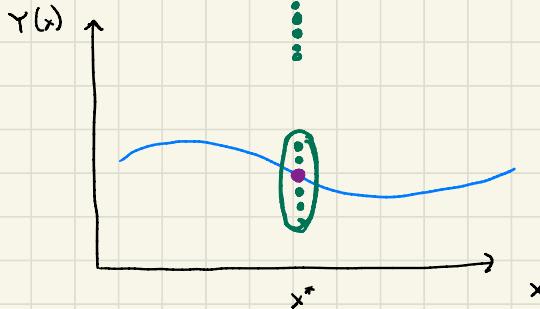
$\hat{Y}(x)$ for different
training
data
 $g(x^*)$
 $g(x)$

High Bias



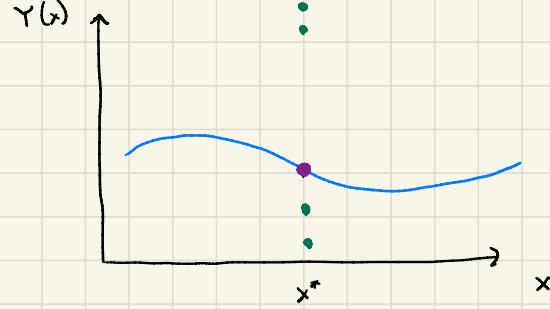
$\hat{Y}(x)$ for different
training
data
 $g(x^*)$
 $g(x)$

Low Variance



$\hat{Y}(x)$ for different
training
data
 $g(x^*)$
 $g(x)$

High Variance



$\hat{Y}(x)$ for different
training
data
 $g(x^*)$
 $g(x)$

We want models with low model bias & low model variance

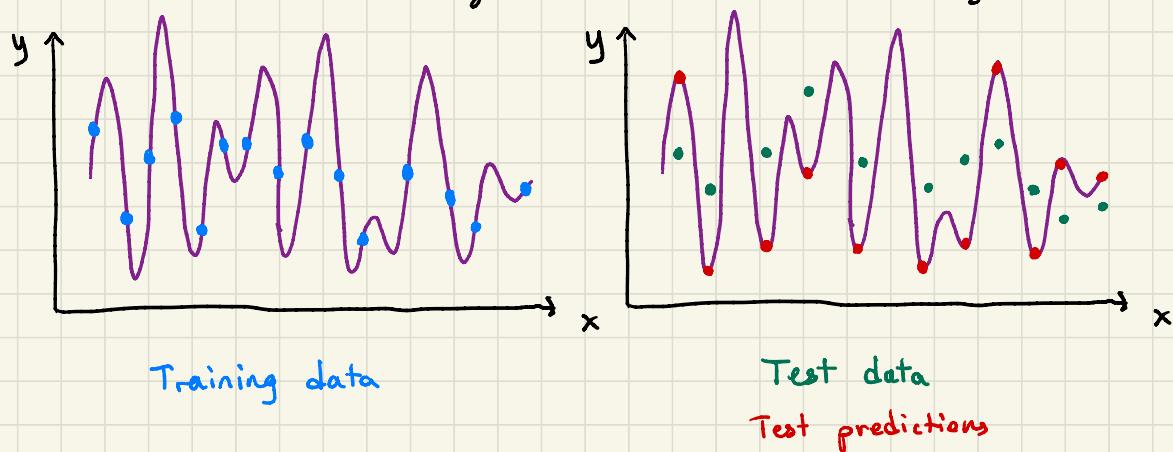
low training error

low validation error

Regularization

high model complexity

- Motivation: If we have a lot of features, we might get a model that performs well on the training data, but does not generalize well (low bias, high variance)



This situation happens because we are giving our model too much freedom \Rightarrow it is using that freedom to find patterns that exist only in the training data

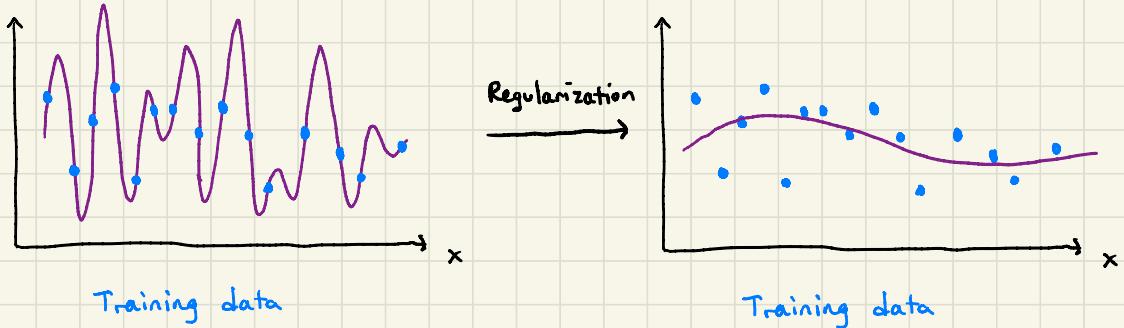
Need to restrict our model to disoverge using all features

$$\text{Before: } \underset{\Theta}{\text{minimize}} \text{ MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - f_{\Theta}(x_i))^2$$

$$\text{Regularization: } \underset{\Theta}{\text{minimize}} \frac{1}{n} \sum_{i=1}^n (y_i - f_{\Theta}(x_i))^2 + \lambda \sum_{j=1}^d \Theta_j^2$$

λ is a scalar that we choose (through cross-validation)

tells the model that choices of Θ with high magnitude are bad!



With regularization, we are sacrificing some training error with the hope that we will get at least as much back on test error (which really matters)

