

Linear Regression

Raguvir Kunani

Data 100, Spring 2020

Simple Linear Regression Quick Review

Simple linear regression involves finding a "line of best fit" that explains the relationship between **2** variables x and y . In fancy math terms:

$$a^*, b^* = \arg \min_{a,b} \sum_{i=1}^n (y_i - (\underbrace{a + bx_i}_{\text{predicted } y \text{ for input } x_i}))^2$$

MSE

sum →
over all n training
data points

All this really means is we are trying to find the values of a and b so that a line with the equation $\hat{y} = a + bx$ best fits the data.

$x^* = \arg \min_x f(x)$ means x^* is equal to the value of x that minimizes $f(x)$

Residuals

The **residual** is a measure of how "off" our prediction was. In simple linear regression, the residual for the i^{th} point is:

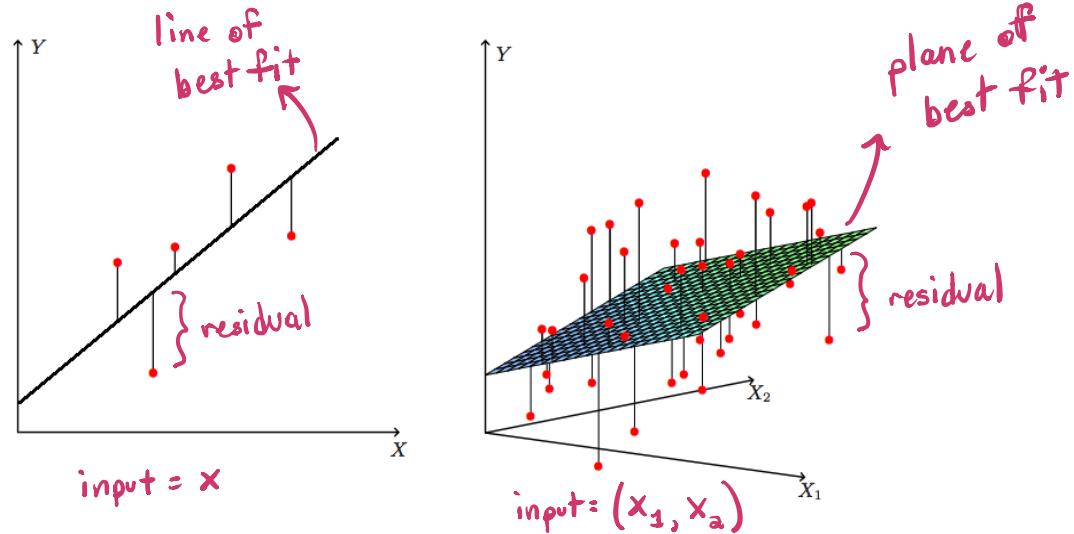
$$\text{true value} \rightarrow y_i - (\underbrace{a + bx_i}_{\text{predicted value}})$$

So really what linear regression is doing is finding the a and b that minimize the sum of (squared) residuals. This view of linear regression as minimizing residuals is very helpful to understanding linear regression in higher dimensions.

Multiple

~~Multivariate~~ Linear Regression in a Graph

In multiple dimensions, linear regression goes from finding a *line* of best fit to finding a *plane* of best fit.



Keep in mind that the quantity you are trying to predict is still a scalar quantity, but now you have multiple inputs (labeled x_1, x_2, \dots, x_d).

Multiple ~~Multivariate~~ Linear Regression in Math

Remember that the simple linear regression problem minimizes the sum of the squared residuals $y_i - (a + bx_i)^2$. We can use the same residuals interpretation to formulate the ~~multivariate~~ linear regression problem, but we need to rewrite the residuals for the multidimensional case.

Let's say our model is of the form $\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2$

Then the residuals become $y_i - (\underbrace{\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2}}_{\text{predicted value}})$

$x_i = \begin{bmatrix} x_{i,1} \\ x_{i,2} \end{bmatrix}$

The sum of the squared residuals is $\sum_{i=1}^n y_i - (\underbrace{\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2}}_{\text{residual}})^2$

$\sum_{\text{over all } n \text{ training data points}}$

MSE

Multiple

~~Multivariate~~ Linear Regression V1

Now we can write our new ~~multiple~~ linear regression problem:

$$\theta_0^*, \theta_1^*, \theta_2^* = \arg \min_{\theta_0, \theta_1, \theta_2} \sum_{i=1}^n (y_i - (\underbrace{\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2}}_{\text{gets uglier as } \# \text{ dimensions increases}}))^2$$

This looks horrendous, and will only get worse with more dimensions.

This is where linear algebra comes into play. Let's rewrite the above problem using matrix vector notation.

Simplifying the Expression for Model's Prediction

Let's start by collecting all the parameters into a vector $\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix}$.

Then we can write $\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2}$ as $\begin{bmatrix} 1 \\ x_{i,1} \\ x_{i,2} \end{bmatrix} \cdot \theta$.

If we define a vector $\phi_i = \begin{bmatrix} 1 \\ x_{i,1} \\ x_{i,2} \end{bmatrix}$,

then we can write $\underbrace{\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2}}$ even more simply as $\phi_i \cdot \theta$.

$$\begin{bmatrix} a \\ b \end{bmatrix} \cdot \begin{bmatrix} c \\ d \end{bmatrix} = ac + bd$$

model's prediction for input x_i

M^{multiple} ~~M~~ultivariate Linear Regression V2

We can update our ~~m~~^{multiple}
~~u~~ltivariate linear regression problem to be:

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^n (y_i - (\phi_i \cdot \theta))^2$$

This looks a lot better, but we can simplify even more.

Let's write the true values and our predicted values as n -dimensional

vectors: $y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$ and $\hat{y} = \begin{bmatrix} \phi_1 \cdot \theta \\ \phi_2 \cdot \theta \\ \vdots \\ \phi_n \cdot \theta \end{bmatrix}$.

\hat{y} is a common notation
for predicted values

n -dimensional because
there are n data points

~~Multiple~~
~~Multivariate~~ Linear Regression V3

Recall $\|z\|$ is the norm of the vector z and $\|z\|^2 = \sum_{i=1}^n z_i^2$.

Then we rewrite our linear regression problem as:

$$\theta^* = \arg \min_{\theta} \|y - \hat{y}\|^2$$

There's just one last simplification. Note that

$$\hat{y} = \begin{bmatrix} \phi_1 \cdot \theta \\ \phi_2 \cdot \theta \\ \vdots \\ \phi_n \cdot \theta \end{bmatrix} = \begin{bmatrix} \text{---} & \phi_1 & \text{---} \\ \text{---} & \phi_2 & \text{---} \\ & \vdots & \\ \text{---} & \phi_n & \text{---} \end{bmatrix} \cdot \theta = \Phi \cdot \theta$$

where Φ is a matrix that contains the ϕ_i as its rows.

Multiple

Multivariate Linear Regression V4

Using Φ , we can write the ~~multiple~~ linear regression problem as:

$$\theta^* = \arg \min_{\theta} \|y - \Phi \cdot \theta\|^2 = \boxed{\arg \min_{\theta} \|y - \Phi \theta\|^2}$$

This is the final version of the problem! This representation is much more compact and *does not change* as the number of dimensions increase.

Note that this version of the problem is *exactly equivalent* to the original version we started with. I did not change the problem; I just rewrote the problem into a form that allows us to use linear algebra.

A Closer Look at $\|y - \Phi\theta\|^2$

Since $\|z\|^2$ represents the length of z , we can think of $\|x - z\|^2$ as the "length of the difference" between x and z .

Put more simply, $\|x - z\|^2$ is the distance between x and z .

This means linear regression, in trying to minimize $\|y - \Phi\theta\|^2$, is really trying to minimize the distance between y and $\Phi\theta$.

This sounds very similar to minimizing residuals! In this case, the residual is a vector, specifically $y - \Phi\theta$.

How to Compute θ^* : Part 1

Since our goal is to choose θ^* such that $\Phi\theta^*$ is as close to y as possible, it seems natural to set $\Phi\theta^* = y$ and solve for θ^* .

However, this equation rarely has a unique solution. This could happen for 2 main reasons:

1. $y \notin \text{range}(\Phi)$. This means there is no vector v such that $\Phi v = y$. (More on this in the next slide).
2. Φ is not invertible. If this is the case, create a new Φ and start over.

Thus, the best we can do is choose θ^* that minimizes the distance between y and $\Phi\theta^*$.

A Closer Look at $y \notin \text{range}(\Phi)$

The matrix-vector product $\Phi\theta$ can be thought of as a linear combination of the columns of Φ .

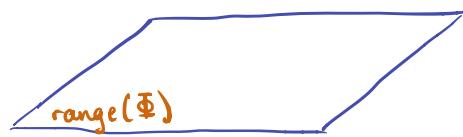
$$\Phi\theta = \theta_0 \Phi_0 + \theta_1 \Phi_1 + \dots + \theta_d \Phi_d$$

Φ_i is the i^{th} column of Φ matrix

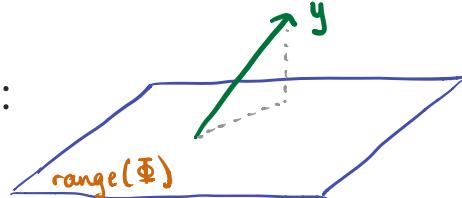
θ_i is the i^{th} element of θ vector

The set of all possible linear combinations of the columns of Φ is called $\text{range}(\Phi)$ because it contains all possible outputs of multiplying Φ by any vector.

Graphically, $\text{range}(\Phi)$ is depicted as:



Then $y \notin \text{range}(\Phi)$ is depicted as:

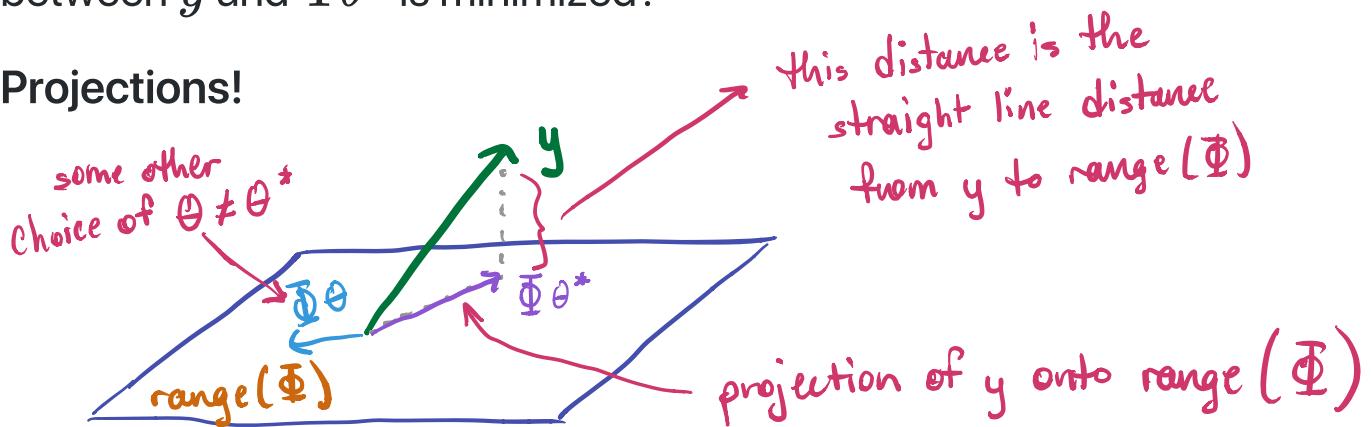


How to Compute θ^* : Part 2

Since we are stuck with $y \notin \text{range}(\Phi)$, we can't set $\Phi\theta^* = y$.

So how else are we supposed to choose θ^* such that the distance between y and $\Phi\theta^*$ is minimized?

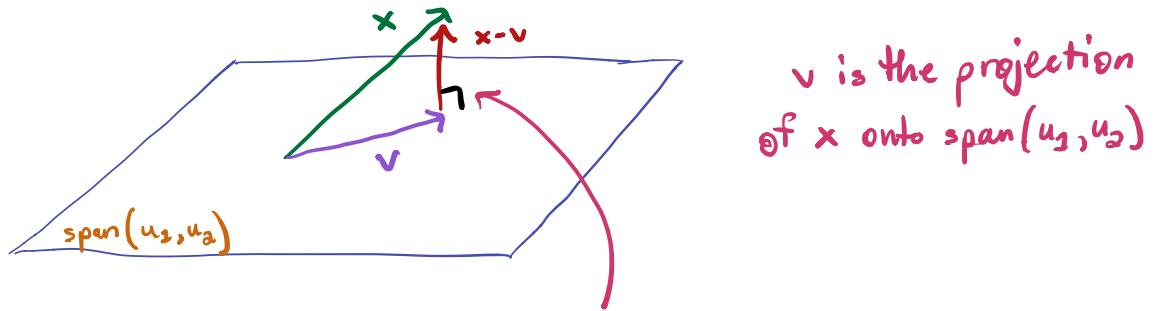
Projections!



You can see that out of all vectors in $\text{range}(\Phi)$, $\Phi\theta^*$ has the closest distance to y

Projections and Their Important Properties

The projection of a vector x onto $\text{span}(u_1, u_2)$ is the vector in the span of u_1 and u_2 that is closest in distance to x .



If v is the projection of x onto $\text{span}(u_1, u_2)$, then $x - v$ is orthogonal to $\text{span}(u_1, u_2)$. This is because the straight line distance is the closest distance between any 2 points.

Using Projections to Solve Linear Regression

From the previous slide, we know that $y - \Phi\theta^*$ is orthogonal to $\text{range}(\Phi)$. (*because $\Phi\theta^*$ is the projection of y onto $\text{range}(\Phi)$*)

Since $\text{range}(\Phi)$ is the set of all linear combinations of the columns of Φ , it must also be true that $y - \Phi\theta^*$ is orthogonal to each column of Φ :

$$\Phi_1 \cdot (y - \Phi\theta^*) = \Phi_2 \cdot (y - \Phi\theta^*) = \dots = \Phi_d \cdot (y - \Phi\theta^*) = 0.$$

d columns of Φ

We can write all these equations using matrix form as:

$$\begin{bmatrix} \Phi_1 \\ \Phi_2 \\ \vdots \\ \Phi_d \end{bmatrix} \cdot (y - \Phi\theta^*) = \Phi^T(y - \Phi\theta^*) = 0$$

Φ^T because Φ_i is a column of Φ

The Normal Equations

Solving the equation from last slide:

$$\Phi^T(y - \Phi\theta^*) = 0$$

$$\Phi^T y - (\Phi^T \Phi)\theta^* = 0$$

$$(\Phi^T \Phi)\theta^* = \Phi^T y$$

These are the normal equations you saw in lecture! Solving for θ^* ,

$$\theta^* = (\Phi^T \Phi)^{-1} \Phi^T y$$

Some Terminology

There are a lot of different names for the process of finding θ^* :

1. Training a least squares model
2. Fitting a least squares model
3. Finding θ that minimizes MSE

All of these things just mean setting up and solving the normal equations $(\Phi^T \Phi) \theta^* = \Phi^T y$.

When you call `model.fit(X, y)`, solving the normal equations is what goes on behind the scenes.

What Happens After Finding θ^* ?

After you've computed θ^* , you can use your model to make predictions!

new data point ϕ vector prediction for x

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow \phi = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow \theta^* \cdot \phi = \hat{y}$$

(1 for bias term)

When you call `model.predict(x)`, this is what is happening behind the scenes.

Feedback Form

This *anonymous* form is for me to learn what I can do to ensure you all get the most out of this class (and specifically, discussion and lab). This form will be open all semester, and I'll be checking it regularly. Be as ruthless as you want, I promise my feelings won't get hurt.

Feedback Form: tinyurl.com/raguvirTAfeedback